



News

[Introducing DeepSeek-V3 2024/12/26](#)

Introducing DeepSeek-V3

Biggest leap forward yet

- ⚡ 60 tokens/second (3x faster than V2!)
- 🧑‍🔬 Enhanced capabilities
- 🔧 API compatibility intact
- 🌍 Fully open-source models & papers




Hi, I'm DeepSeek.

How can I help you today?

Write a romance novel

 DeepThink




 Search **NEW**



AI-generated, for reference only

Benchmark (Metric)		DeepSeek-V3	Qwen2.5 72B-Inst.	Llama3.1 405B-Inst.	Claude-3.5- Sonnet-1022	GPT-4o 0513
Architecture		MoE	Dense	Dense	-	-
# Activated Params		37B	72B	405B	-	-
# Total Params		671B	72B	405B	-	-
English	MMLU (EM)	88.5	85.3	88.6	88.3	87.2
	MMLU-Redux (EM)	89.1	85.6	86.2	88.9	88
	MMLU-Pro (EM)	75.9	71.6	73.3	78	72.6
	DROP (3-shot F1)	91.6	76.7	88.7	88.3	83.7
	IF-Eval (Prompt Strict)	86.1	84.1	86	86.5	84.3
	GPQA-Diamond (Pass@1)	59.1	49	51.1	65	49.9
	SimpleQA (Correct)	24.9	9.1	17.1	28.4	38.2
	FRAMES (Acc.)	73.3	69.8	70	72.5	80.5
	LongBench v2 (Acc.)	48.7	39.4	36.1	41	48.1
Code	HumanEval-Mul (Pass@1)	82.6	77.3	77.2	81.7	80.5
	LiveCodeBench(Pass@1-COT)	40.5	31.1	28.4	36.3	33.4
	LiveCodeBench (Pass@1)	37.6	28.7	30.1	32.8	34.2
	Codeforces (Percentile)	51.6	24.8	25.3	20.3	23.6
	SWE Verified (Resolved)	42	23.8	24.5	50.8	38.8
	Aider-Edit (Acc.)	79.7	65.4	63.9	84.2	72.9
	Aider-Polyglot (Acc.)	49.6	7.6	5.8	45.3	16
Math	AIME 2024 (Pass@1)	39.2	23.3	23.3	16	9.3
	MATH-500 (EM)	90.2	80	73.8	78.3	74.6
	CNMO 2024 (Pass@1)	43.2	15.9	6.8	13.1	10.8
Chinese	CLUEWSC (EM)	90.9	91.4	84.7	85.4	87.9
	C-Eval (EM)	86.5	86.1	61.5	76.7	76
	C-SimpleQA (Correct)	64.1	48.4	50.4	51.3	59.3



What's new in V3

-  671B MoE parameters
-  37B activated parameters
-  Trained on 14.8T high-quality tokens

 Dive deeper here:

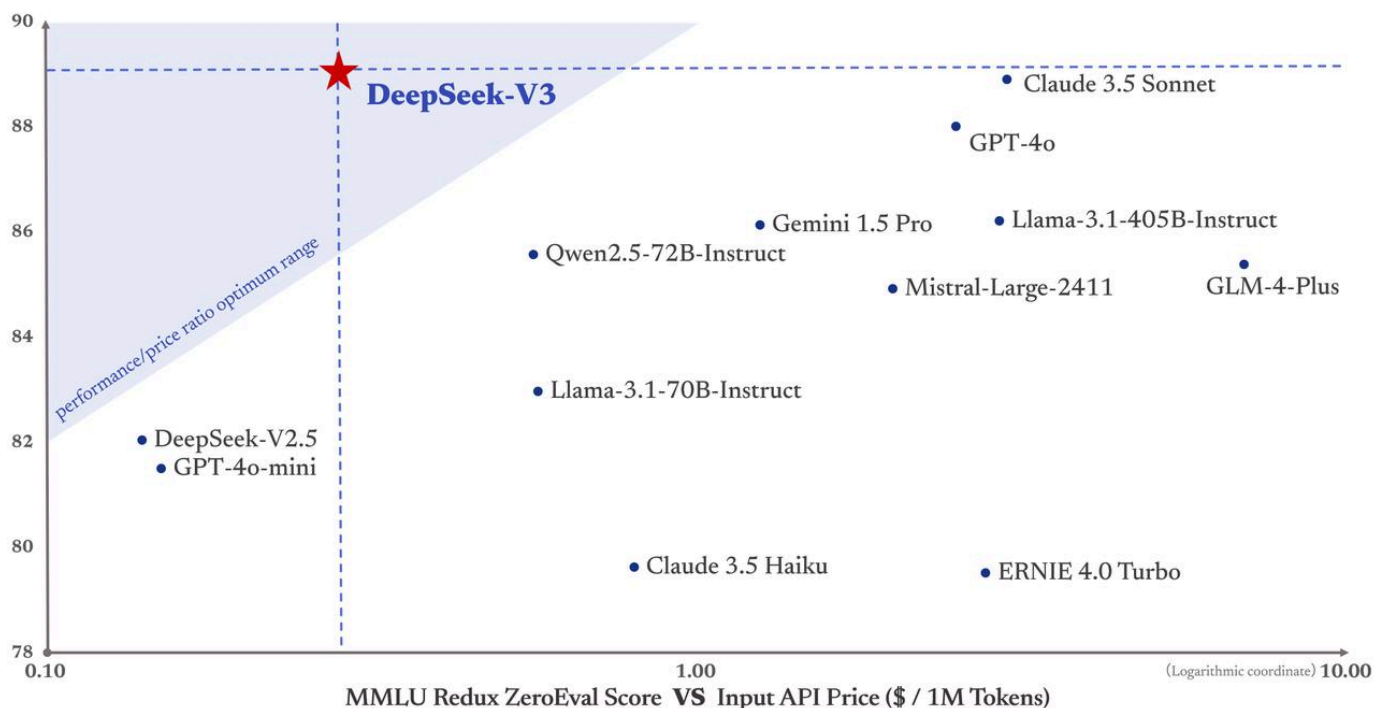
- Model  <https://github.com/deepseek-ai/DeepSeek-V3>
- Paper  https://github.com/deepseek-ai/DeepSeek-V3/blob/main/DeepSeek_V3.pdf

API Pricing Update

-  Until Feb 8: same as V2!
-  From Feb 8 onwards:

Input (cache miss)	Input (cache hit)	Output
\$0.27/M tokens	\$0.07/M tokens	\$1.10/M tokens

Still the best value in the market! 🔥



DeepSeek-V3 API

Pay-as-you-go pricing

Input :

cache hit

~~\$0.07~~ **\$0.014** / 1M tokens

cache miss

~~\$0.27~~ **\$0.14** / 1M tokens



Output :

~~\$1.10~~ **\$0.28** / 1M tokens



The promotional period will end
after February 8, 2025, at 16:00 UTC

- 🚀 Open-source spirit + Longtermism to inclusive AGI
- ☀️ DeepSeek's mission is unwavering. We're thrilled to share our progress with the community and see the gap between open and closed models narrowing.

-  This is just the beginning! Look forward to multimodal support and other cutting-edge features in the DeepSeek ecosystem.
-  Together, let's push the boundaries of innovation!